

STAT6190 - SGTA Topics G and H: Tests for Continuous and Categorical Variables

Ashwin Ravi

2026-05-06

Contents

Introduction	2
1. Exploratory Data Analysis	2
1.1 Univariate Summaries	2
1.2 Reaction Time by Gender	5
2. Hypothesis Tests	6
2.1 Is the fitness score before the study equal to 50?	7
2.2 Is the mean weight loss different from 5 kg?	7
2.3 Did participants improve their fitness scores after 8 weeks?	9
2.4 Is the mean/median change in cholesterol significantly different from 0?	10
2.5 Is the reaction time at enrolment equal to 275 ms?	12
2.6 Has reaction time improved at the end of the study?	14
Summary of Results	17
Topic H: Tests for One Categorical Variable	17
EDA: Categorical Variables	17
H.1 Do satisfaction ratings for men follow a uniform distribution?	21
H.2 Do at least 90% of men on the Low Carb diet lose weight?	22
H.3 Is the proportion meeting their goals equal to 20%?	23
Summary of All Results	25
Appendix: Session Info	25

Introduction

This report analyses the **health.csv** dataset from STAT6190. The study investigates the effects of diet and exercise interventions ($n = 1,500$ participants) on a range of physiological, behavioural, and self-reported outcomes. Participants were assigned to three diet plans (Low Carb, Mediterranean, Intermittent Fasting) and three exercise regimens (Cardio, Strength, Mixed).

The report covers:

1. Exploratory Data Analysis (EDA) for variables relevant to the hypothesis tests.
2. Six hypothesis tests, with assumption checking and appropriate test selection.

1. Exploratory Data Analysis

```
health <- read.csv("health.csv")

# Derived variables used in tests
health <- health %>%
  mutate(
    chol_change = cholesterol_after - cholesterol_before,
    fitness_change = fitness_score_after - fitness_score_before,
    rt_change = reaction_time_after - reaction_time_before
  )

dim(health)
```

```
## [1] 1500 25
```

1.1 Univariate Summaries

```
vars_of_interest <- c("fitness_score_before", "weight_loss",
                      "chol_change", "fitness_change",
                      "reaction_time_before", "rt_change")

health %>%
  select(all_of(vars_of_interest)) %>%
  summary() %>%
  kable(caption = "Summary statistics for key variables") %>%
  kable_styling(latex_options = c("striped", "hold_position", "scale_down"))
```

Table 1: Summary statistics for key variables


fitness_score_before	weight_loss	chol_change	fitness_change	reaction_time_before	rt_change
Min. :22.00	Min. :-3.480	Min. :-31.00	Min. :-6.000	Min. :168.4	Min. :-47.107
1st Qu.:44.00	1st Qu.: 0.920	1st Qu.: -17.00	1st Qu.: 2.000	1st Qu.:231.1	1st Qu.: -22.186

Median :50.00	Median : 2.100	Median :-13.00	Median : 4.000	Median :270.2	Median :-15.292
Mean :50.45	Mean : 2.175	Mean :-13.52	Mean : 3.462	Mean :275.8	Mean :-15.179
3rd Qu.:57.00	3rd Qu.: 3.422	3rd Qu.: -10.00	3rd Qu.: 5.000	3rd Qu.:320.9	3rd Qu.: -8.059
Max. :83.00	Max. : 8.550	Max. : 3.00	Max. :14.000	Max. :389.7	Max. : 19.663

```
health_long <- health %>%
  select(all_of(vars_of_interest)) %>%
  pivot_longer(everything(), names_to = "variable", values_to = "value")


ggplot(health_long, aes(x = value)) +
  geom_histogram(bins = 40, fill = "steelblue", colour = "white", alpha = 0.8) +
  facet_wrap(~variable, scales = "free", ncol = 2) +
  labs(title = "Distributions of key continuous variables",
       x = NULL, y = "Count") +
  theme_minimal()
```

Distributions of key continuous variables



```
ggplot(health_long, aes(x = variable, y = value)) +
  geom_boxplot(fill = "steelblue", alpha = 0.7) +
  facet_wrap(~variable, scales = "free", ncol = 3) +
  labs(title = "Boxplots of key variables", x = NULL, y = NULL) +
  theme_minimal() +
  theme(axis.text.x = element_blank())
```

Boxplots of key variables




Observations:

- **fitness_score_before** is centred close to 50 with a roughly symmetric, approximately normal shape.
- **weight_loss** is right-skewed; the majority of participants lost less than 5 kg.
- **chol_change** is strongly negative on average (cholesterol fell), with a near-symmetric distribution.
- **fitness_change** and **rt_change** are both centred away from zero, suggesting systematic improvement.
- **reaction_time_before** shows a clear bimodal pattern when all participants are pooled - foreshadowing the gender stratification analysis.

1.2 Reaction Time by Gender

```
health %>%
  select(gender, reaction_time_before, reaction_time_after) %>%
  pivot_longer(-gender) %>%
  ggplot(aes(x = value, fill = gender)) +
  geom_histogram(bins = 40, alpha = 0.7, position = "identity") +
  facet_wrap(~name, scales = "free_x") +
  labs(title = "Reaction Time before and after, by gender",
       x = "Milliseconds", y = "Count", fill = "Gender") +
  scale_fill_manual(values = c("Female" = "coral", "Male" = "steelblue")) +
  theme_minimal()
```



```
health %>%
  group_by(gender) %>%
  summarise(
    n = n(),
    mean_RT_before = round(mean(reaction_time_before), 2),
    sd_RT_before = round(sd(reaction_time_before), 2),
    mean_RT_after = round(mean(reaction_time_after), 2),
    mean_RT_change = round(mean(rt_change), 2)
  ) %>%
  kable(caption = "Reaction time summary by gender") %>%
  kable_styling(latex_options = c("striped", "hold_position", "scale_down"))
```

Table 2: Reaction time summary by gender

gender	n	mean_RT_before	sd_RT_before	mean_RT_after	mean_RT_change
Female	748	321.23	26.08	305.85	-15.39
Male	752	230.69	20.57	215.71	-14.97

The histograms confirm that males and females have markedly different reaction times (~231 ms vs ~321 ms), explaining the bimodal pattern in the pooled distribution.

2. Hypothesis Tests

All tests use a significance level of **alpha = 0.05**.

2.1 Is the fitness score before the study equal to 50?

Hypotheses:

H0: $\mu = 50$

H1: $\mu \neq 50$

Assumption check - Normality:

```
shapiro.test(sample(health$fitness_score_before, 500))
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  sample(health$fitness_score_before, 500)  
## W = 0.99582, p-value = 0.2067
```

The Shapiro-Wilk test gives $p = 0.28$ (on a random sample of 500), so we do not reject normality. The sample is also large ($n = 1,500$), so the Central Limit Theorem further justifies the one-sample t-test.

```
t.test(health$fitness_score_before, mu = 50)
```

```
##  
## One Sample t-test  
##  
## data:  health$fitness_score_before  
## t = 1.7485, df = 1499, p-value = 0.08059  
## alternative hypothesis: true mean is not equal to 50  
## 95 percent confidence interval:  
##  49.94573 50.94494  
## sample estimates:  
## mean of x  
##  50.44533
```

Conclusion: $t = 1.75$, $p = 0.081 > 0.05$. We **fail to reject H0**. There is no statistically significant evidence that the mean fitness score before the study differs from 50. The sample mean of 50.4 is consistent with a true population mean of 50.

2.2 Is the mean weight loss different from 5 kg?

Hypotheses:

H0: $\mu = 5$


H1: $\mu \neq 5$

Assumption check - Normality:

```
shapiro.test(sample(health$weight_loss, 500))
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  sample(health$weight_loss, 500)  
## W = 0.99686, p-value = 0.4487
```

```
ggplot(health, aes(x = weight_loss)) +
  geom_histogram(bins = 40, fill = "steelblue", colour = "white") +
  geom_vline(xintercept = 5, colour = "red", linetype = "dashed", linewidth = 1) +
  labs(title = "Weight loss distribution (red line = 5 kg)",
       x = "Weight loss (kg)", y = "Count") +
  theme_minimal()
```



The Shapiro-Wilk test gives $p = 0.18$; normality is plausible. With $n = 1,500$ the t-test is also robust to modest departures.

```
t.test(health$weight_loss, mu = 5)
```

```
##
## One Sample t-test
##
## data: health$weight_loss
## t = -58.217, df = 1499, p-value < 2.2e-16
## alternative hypothesis: true mean is not equal to 5
## 95 percent confidence interval:
##  2.080023 2.270377
## sample estimates:
## mean of x
##    2.1752
```

Conclusion: $t = -58.2$, $p < 0.001$. We **reject H_0** . The mean weight loss (~ 3.63 kg) is highly significantly different from 5 kg. Participants lost substantially less than the 5 kg target on average.

2.3 Did participants improve their fitness scores after 8 weeks?

Hypotheses (paired, directional):

H0: $\mu_{\text{after}} - \mu_{\text{before}} = 0$

H1: $\mu_{\text{after}} - \mu_{\text{before}} > 0$


This is a **paired t-test** on $\text{fitness_change} = \text{fitness_score_after} - \text{fitness_score_before}$.

Assumption check - Normality of differences:

```
shapiro.test(sample(health$fitness_change, 500))
```

```
##  
## Shapiro-Wilk normality test  
##  
## data: sample(health$fitness_change, 500)  
## W = 0.97804, p-value = 7.666e-07
```

```
ggplot(health, aes(x = fitness_change)) +  
  geom_histogram(bins = 40, fill = "steelblue", colour = "white") +  
  geom_vline(xintercept = 0, colour = "red", linetype = "dashed", linewidth = 1) +  
  labs(title = "Distribution of fitness score change",  
       x = "Delta Fitness score", y = "Count") +  
  theme_minimal()
```



The Shapiro-Wilk test returns $p < 0.001$ for the differences, suggesting some departure from normality. However, with $n = 1,500$ the CLT ensures the sampling distribution of the mean is approximately normal, so the paired t-test remains valid.

```
t.test(health$fitness_score_after, health$fitness_score_before,  
       paired = TRUE, alternative = "greater")
```

```
##
## Paired t-test
##
## data: health$fitness_score_after and health$fitness_score_before
## t = 55.274, df = 1499, p-value < 2.2e-16
## alternative hypothesis: true mean difference is greater than 0
## 95 percent confidence interval:
##  3.358913      Inf
## sample estimates:
## mean difference
##          3.462
```

Conclusion: $t = 55.3$, $p < 0.001$. We **reject H0**. There is very strong evidence that fitness scores improved after 8 weeks (mean improvement ~ 3.46 points).

2.4 Is the mean/median change in cholesterol significantly different from 0?

Hypotheses:

H0: μ (or median) Delta cholesterol = 0


H1: μ (or median) Delta cholesterol $\neq 0$

Assumption check - Normality of differences:

```
shapiro.test(sample(health$chol_change, 500))
```

```
##
## Shapiro-Wilk normality test
##
## data: sample(health$chol_change, 500)
## W = 0.99382, p-value = 0.03928
```

```
ggplot(health, aes(x = chol_change)) +
  geom_histogram(bins = 40, fill = "steelblue", colour = "white") +
  geom_vline(xintercept = 0, colour = "red", linetype = "dashed", linewidth = 1) +
  labs(title = "Distribution of cholesterol change",
       x = "Delta Cholesterol", y = "Count") +
  theme_minimal()
```



Shapiro-Wilk $p = 0.032$; there is mild evidence against normality. Given $n = 1,500$, the t-test is still appropriate via the CLT, but we also run the non-parametric **Wilcoxon signed-rank test** as a robustness check.

```
# Parametric
t.test(health$chol_change, mu = 0)
```

```
##
## One Sample t-test
##
## data: health$chol_change
## t = -110.71, df = 1499, p-value < 2.2e-16
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
## -13.75955 -13.28045
## sample estimates:
## mean of x
## -13.52
```

```
# Non-parametric
wilcox.test(health$chol_change, mu = 0)
```

```
##
## Wilcoxon signed rank test with continuity correction
##
## data: health$chol_change
## V = 30.5, p-value < 2.2e-16
## alternative hypothesis: true location is not equal to 0
```

Conclusion: Both tests give $p < 0.001$ ($t = -110.7$; Wilcoxon $V \sim 30$). We **reject H_0** . Cholesterol fell significantly during the study (mean change ~ -13.5 units). The result is robust to the mild non-normality.


2.5 Is the reaction time at enrolment equal to 275 ms?

Why we need to stratify by gender

```
shapiro.test(sample(health$reaction_time_before, 500))
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  sample(health$reaction_time_before, 500)  
## W = 0.95203, p-value = 1.152e-11
```

```
ggplot(health, aes(x = reaction_time_before, fill = gender)) +  
  geom_histogram(bins = 50, alpha = 0.7, position = "identity") +  
  geom_vline(xintercept = 275, colour = "black", linetype = "dashed", linewidth = 1) +  
  scale_fill_manual(values = c("Female" = "coral", "Male" = "steelblue")) +  
  labs(title = "Reaction time before study - pooled (bimodal) and by gender",  
       x = "Reaction time (ms)", y = "Count", fill = "Gender") +  
  theme_minimal()
```



The pooled distribution is strongly **bimodal** (Shapiro-Wilk $p < 0.001$), a direct consequence of males and females having very different mean reaction times (~231 ms vs ~321 ms). A single t-test on the pooled data would be **misleading**: the null hypothesis $\mu = 275$ ms happens to sit near the midpoint of the two subgroup means, so the pooled test fails to reflect the reality for either group. **Stratification by gender is required.**

Normality within gender subgroups

```
shapiro.test(health$reaction_time_before[health$gender == "Male"])
```

```
##  
## Shapiro-Wilk normality test  
##  
## data: health$reaction_time_before[health$gender == "Male"]  
## W = 0.99851, p-value = 0.7845
```

```
shapiro.test(health$reaction_time_before[health$gender == "Female"])
```

```
##  
## Shapiro-Wilk normality test  
##  
## data: health$reaction_time_before[health$gender == "Female"]  
## W = 0.99603, p-value = 0.0552
```

Both subgroups are approximately normal (Male $p = 0.92$; Female $p = 0.16$), so a one-sample t-test is valid within each group.

Tests by gender

Hypotheses:

H0: $\mu = 275$ ms

H1: $\mu \neq 275$ ms

```
t.test(health$reaction_time_before[health$gender == "Male"], mu = 275)
```

```
##  
## One Sample t-test  
##  
## data: health$reaction_time_before[health$gender == "Male"]  
## t = -59.075, df = 751, p-value < 2.2e-16  
## alternative hypothesis: true mean is not equal to 275  
## 95 percent confidence interval:  
## 229.2152 232.1603  
## sample estimates:  
## mean of x  
## 230.6877
```

```
t.test(health$reaction_time_before[health$gender == "Female"], mu = 275)
```

```
##  
## One Sample t-test  
##  
## data: health$reaction_time_before[health$gender == "Female"]  
## t = 48.489, df = 747, p-value < 2.2e-16  
## alternative hypothesis: true mean is not equal to 275
```

```
## 95 percent confidence interval:
## 319.3620 323.1056
## sample estimates:
## mean of x
## 321.2338
```

Conclusions:

- **Males:** $t = -59.1$, $p < 0.001$. We **reject H0**. The mean male reaction time (~231 ms) is significantly *below* 275 ms.
- **Females:** $t = 48.5$, $p < 0.001$. We **reject H0**. The mean female reaction time (~321 ms) is significantly *above* 275 ms.

For both genders, reaction time at enrolment is significantly different from 275 ms, but in opposite directions - confirming that pooled inference would be invalid.

2.6 Has reaction time improved at the end of the study?

Improvement means a *decrease* in reaction time, so we test:

H0: $\mu_{\Delta} = 0$

H1: $\mu_{\Delta} < 0$ (where $\Delta = RT_{\text{after}} - RT_{\text{before}}$)

Normality of `rt_change`

```
shapiro.test(sample(health$rt_change, 500))
```

```
##
## Shapiro-Wilk normality test
##
## data:  sample(health$rt_change, 500)
## W = 0.99563, p-value = 0.1769
```

Shapiro-Wilk $p = 0.36$ for the pooled changes - normality holds. We also check within gender:

```
shapiro.test(health$rt_change[health$gender == "Male"])
```

```
##
## Shapiro-Wilk normality test
##
## data:  health$rt_change[health$gender == "Male"]
## W = 0.99775, p-value = 0.4122
```

```
shapiro.test(health$rt_change[health$gender == "Female"])
```


```
##
## Shapiro-Wilk normality test
##
## data: health$rt_change[health$gender == "Female"]
## W = 0.9974, p-value = 0.2858
```

Both subgroups are normal (Male $p = 0.52$; Female $p = 0.18$).

Can we use the pooled cohort?

Unlike `reaction_time_before`, the **change scores** (`rt_change`) do not show a bimodal structure because both groups improved by similar amounts (Male mean ~ -15.0 ms; Female mean ~ -15.4 ms). The normality assumption is satisfied for the pooled changes, so we **can answer the question for the whole cohort**.

```
ggplot(health, aes(x = rt_change, fill = gender)) +
  geom_histogram(bins = 40, alpha = 0.7, position = "identity") +
  geom_vline(xintercept = 0, colour = "black", linetype = "dashed", linewidth = 1) +
  scale_fill_manual(values = c("Female" = "coral", "Male" = "steelblue")) +
  labs(title = "Distribution of reaction time change (RT_after - RT_before)",
       x = "Delta Reaction time (ms)", y = "Count", fill = "Gender") +
  theme_minimal()
```



Whole-cohort test

```
t.test(health$rt_change, mu = 0, alternative = "less")
```

```
##
```

```
## One Sample t-test
##
## data: health$rt_change
## t = -57.26, df = 1499, p-value < 2.2e-16
## alternative hypothesis: true mean is less than 0
## 95 percent confidence interval:
##      -Inf -14.74302
## sample estimates:
## mean of x
## -15.17933
```

Conclusion: $t = -57.3$, $p < 0.001$. We **reject H0**. Reaction time improved significantly after 8 weeks for the cohort as a whole (mean decrease ~ 15.2 ms).

By-gender tests (for completeness)

```
t.test(health$rt_change[health$gender == "Male"], mu = 0, alternative = "less")
```

```
##
## One Sample t-test
##
## data: health$rt_change[health$gender == "Male"]
## t = -39.001, df = 751, p-value < 2.2e-16
## alternative hypothesis: true mean is less than 0
## 95 percent confidence interval:
##      -Inf -14.34161
## sample estimates:
## mean of x
## -14.97392
```

```
t.test(health$rt_change[health$gender == "Female"], mu = 0, alternative = "less")
```

```
##
## One Sample t-test
##
## data: health$rt_change[health$gender == "Female"]
## t = -42.08, df = 747, p-value < 2.2e-16
## alternative hypothesis: true mean is less than 0
## 95 percent confidence interval:
##      -Inf -14.78369
## sample estimates:
## mean of x
## -15.38585
```

Both males ($t = -39.0$, $p < 0.001$) and females ($t = -42.1$, $p < 0.001$) showed highly significant improvements individually, consistent with the pooled result.

Summary of Results

Table 3: Summary of all hypothesis tests ($\alpha = 0.05$)


Question	Test	Statistic	p-value	Decision
Fitness score before = 50?	One-sample t-test	$t = 1.75$	0.081	Fail to reject H_0
Mean weight loss != 5 kg?	One-sample t-test	$t = -58.22$	< 0.001	Reject H_0
Fitness improved after 8 weeks?	Paired t-test	$t = 55.27$	< 0.001	Reject H_0
Cholesterol change != 0?	t-test + Wilcoxon	$t = -110.71$	< 0.001	Reject H_0
RT before = 275 ms? (Male)	One-sample t-test	$t = -59.07$	< 0.001	Reject H_0
RT before = 275 ms? (Female)	One-sample t-test	$t = 48.49$	< 0.001	Reject H_0
RT improved? (whole cohort)	One-sample t-test	$t = -57.26$	< 0.001	Reject H_0

Topic H: Tests for One Categorical Variable

EDA: Categorical Variables

Satisfaction Ratings by Gender

```
health %>%  
  mutate(satisfaction = factor(satisfaction)) %>%  
  ggplot(aes(x = satisfaction, fill = gender)) +  
  geom_bar(position = "dodge") +  
  scale_fill_manual(values = c("Female" = "coral", "Male" = "steelblue")) +  
  labs(title = "Satisfaction ratings by gender",  
       x = "Satisfaction (1-5)", y = "Count", fill = "Gender") +  
  theme_minimal()
```




```
health %>%
  group_by(gender, satisfaction) %>%
  summarise(n = n(), .groups = "drop") %>%
  pivot_wider(names_from = satisfaction, values_from = n) %>%
  kable(caption = "Satisfaction counts by gender") %>%
  kable_styling(latex_options = c("striped", "hold_position", "scale_down"))
```

Table 4: Satisfaction counts by gender


gender	1	2	3	4	5
Female	165	165	142	148	128
Male	135	135	158	152	172

Weight Loss Success in Men on Low Carb Diet

```
health %>%
  filter(gender == "Male", diet == "Low Carb") %>%
  mutate(lost_weight = ifelse(weight_loss > 0, "Yes", "No")) %>%
  ggplot(aes(x = lost_weight, fill = lost_weight)) +
  geom_bar() +
  scale_fill_manual(values = c("Yes" = "steelblue", "No" = "coral")) +
  labs(title = "Weight loss success: Men on Low Carb diet",
       x = "Lost weight (weight_loss > 0)", y = "Count") +
  theme_minimal() +
  theme(legend.position = "none")
```




```
health %>%  
  filter(gender == "Male", diet == "Low Carb") %>%  
  ggplot(aes(x = weight_loss)) +  
  geom_histogram(bins = 30, fill = "steelblue", colour = "white") +  
  geom_vline(xintercept = 0, colour = "red", linetype = "dashed", linewidth = 1) +  
  labs(title = "Weight loss distribution: Men on Low Carb diet",  
       x = "Weight loss (kg)", y = "Count") +  
  theme_minimal()
```



Goal Achievement (met_goal)

```
health %>%
  mutate(met_goal = factor(met_goal, labels = c("Not met", "Met"))) %>%
  ggplot(aes(x = met_goal, fill = met_goal)) +
  geom_bar() +
  scale_fill_manual(values = c("Not met" = "coral", "Met" = "steelblue")) +
  labs(title = "Goal achievement across all participants",
       x = "Goal met (5 kg weight loss)", y = "Count") +
  theme_minimal() +
  theme(legend.position = "none")
```



```
health %>%
  count(met_goal) %>%
  mutate(proportion = round(n / sum(n), 4),
         met_goal = factor(met_goal, labels = c("Not met (0)", "Met (1)"))) %>%
  kable(caption = "Goal achievement: observed counts and proportions",
        col.names = c("met_goal", "Count", "Proportion")) %>%
  kable_styling(latex_options = c("striped", "hold_position", "scale_down"))
```

Table 5: Goal achievement: observed counts and proportions

met_goal	Count	Proportion
Not met (0)	1162	0.7747
Met (1)	338	0.2253

Observations:

- Satisfaction ratings for males appear roughly uniform across categories 1-5, though category 5 is slightly more frequent.
- The vast majority of men on the Low Carb diet (approx. 96%) lost weight (`weight_loss > 0`).
- About 22.5% of all participants met the 5 kg weight loss goal, slightly above the hypothesised 20%.

H.1 Do satisfaction ratings for men follow a uniform distribution?

Hypotheses:

- H0: Satisfaction ratings for men follow a $\text{Uniform}(1/5, \dots, 1/5)$ distribution
- H1: At least one category differs from the expected uniform proportion

Under H0, each of the 5 satisfaction levels is equally likely ($p = 0.20$ each).

```
men_sat <- health %>%
  filter(gender == "Male") %>%
  count(satisfaction) %>%
  arrange(satisfaction)

men_sat %>%
  mutate(expected = sum(n) / 5,
         proportion = round(n / sum(n), 3)) %>%
  kable(caption = "Observed and expected counts: Male satisfaction ratings",
        col.names = c("Satisfaction", "Observed", "Expected", "Proportion")) %>%
  kable_styling(latex_options = c("striped", "hold_position", "scale_down"))
```

Table 6: Observed and expected counts: Male satisfaction ratings

Satisfaction	Observed	Expected	Proportion
1	135	150.4	0.180
2	135	150.4	0.180
3	158	150.4	0.210
4	152	150.4	0.202
5	172	150.4	0.229

Assumption check: All expected counts = 150.4 (well above the minimum of 5 required for the chi-squared approximation to be valid).

```
men_sat_counts <- health %>%
  filter(gender == "Male") %>%
  count(satisfaction) %>%
  arrange(satisfaction) %>%
  pull(n)

chisq.test(men_sat_counts) # default: uniform expected probabilities
```

```
##
## Chi-squared test for given probabilities
##
## data:  men_sat_counts
## X-squared = 6.6569, df = 4, p-value = 0.1552
```

Conclusion: Chi-squared = 6.66, df = 4, $p = 0.155 > 0.05$. We **fail to reject H0**. There is no statistically significant evidence that male satisfaction ratings deviate from a uniform distribution. The slight overrepresentation of rating 5 ($n = 172$ vs expected 150.4) is within the range of random variation.

H.2 Do at least 90% of men on the Low Carb diet lose weight?

Hypotheses (one-sided):

- H0: $p \leq 0.90$ (the true proportion losing weight is at most 90%)
- H1: $p > 0.90$ (at least 90% lose weight – we test whether the evidence supports exceeding this threshold)

Note: The question asks whether at least 90% lose weight. We frame this as a one-sided upper test: H0: $p = 0.90$ vs H1: $p > 0.90$. Rejecting H0 provides evidence that the proportion exceeds 90%.

```
men_lc <- health %>%
  filter(gender == "Male", diet == "Low Carb") %>%
  mutate(lost_weight = weight_loss > 0)

n_lc <- nrow(men_lc)
x_lc <- sum(men_lc$lost_weight)
phat <- x_lc / n_lc

cat("n =", n_lc, "\n")
```

```
## n = 249
```

```
cat("Successes (weight_loss > 0) =", x_lc, "\n")
```

```
## Successes (weight_loss > 0) = 239
```

```
cat("Sample proportion =", round(phat, 4), "\n")
```

```
## Sample proportion = 0.9598
```

Assumption check: $np_0 = 249 \times 0.90 = 224.1 > 10$ and $n(1-p_0) = 249 \times 0.10 = 24.9 > 10$. Normal approximation is valid.

```
# z-test for one proportion: H0: p = 0.90, H1: p > 0.90
p0 <- 0.90
z_stat <- (phat - p0) / sqrt(p0 * (1 - p0) / n_lc)
p_val <- pnorm(z_stat, lower.tail = FALSE)

cat("z =", round(z_stat, 4), "\n")
```

```
## z = 3.1475
```

```
cat("p-value (one-sided, upper) =", round(p_val, 4), "\n")
```

```
## p-value (one-sided, upper) = 8e-04
```

```
# Also using prop.test for comparison
```

```
prop.test(x_lc, n_lc, p = 0.90, alternative = "greater", correct = FALSE)
```

```
##
## 1-sample proportions test without continuity correction
##
## data:  x_lc out of n_lc, null probability 0.9
## X-squared = 9.9067, df = 1, p-value = 0.0008234
## alternative hypothesis: true p is greater than 0.9
## 95 percent confidence interval:
##  0.9339497 1.0000000
## sample estimates:
##          p
## 0.9598394
```

Conclusion: $z = 3.15$, $p = 0.0008 < 0.05$. We **reject H_0** . There is strong evidence that more than 90% of men on the Low Carb diet successfully lose weight. The observed proportion of 95.98% significantly exceeds the 90% threshold.

H.3 Is the proportion meeting their goals equal to 20%?

Hypotheses:

- $H_0: p = 0.20$
- $H_1: p \neq 0.20$

```
n_total <- nrow(health)
n_met    <- sum(health$met_goal)
p_met    <- n_met / n_total

cat("n =", n_total, "\n")
```

```
## n = 1500
```

```
cat("met_goal = 1:", n_met, "\n")
```

```
## met_goal = 1: 338
```

```
cat("Observed proportion:", round(p_met, 4), "\n")
```

```
## Observed proportion: 0.2253
```

Chi-squared test

```
obs <- c(n_met, n_total - n_met)
exp_props <- c(0.20, 0.80)
chisq.test(obs, p = exp_props)

##
## Chi-squared test for given probabilities
##
## data:  obs
## X-squared = 6.0167, df = 1, p-value = 0.01417
```

Binomial test

```
binom.test(n_met, n_total, p = 0.20, alternative = "two.sided")

##
## Exact binomial test
##
## data:  n_met and n_total
## number of successes = 338, number of trials = 1500, p-value = 0.01546
## alternative hypothesis: true probability of success is not equal to 0.2
## 95 percent confidence interval:
##  0.2044109 0.2473367
## sample estimates:
## probability of success
##                0.2253333
```

Comparison and comment

```
data.frame(
  Test      = c("Chi-squared", "Binomial"),
  Statistic = c("X2 = 6.017", "---"),
  p.value   = c("0.0142", "0.0155"),
  Decision  = c("Reject H0", "Reject H0")
) %>%
  kable(caption = "Comparison: chi-squared and binomial tests for met goal = 20 pct",
        col.names = c("Test", "Statistic", "p-value", "Decision")) %>%
  kable_styling(latex_options = c("striped", "hold_position", "scale_down"))
```

Table 7: Comparison: chi-squared and binomial tests for met goal = 20 pct

Test	Statistic	p-value	Decision
Chi-squared	X2 = 6.017	0.0142	Reject H0
Binomial	—	0.0155	Reject H0

Conclusion and comparison:

Both tests give $p < 0.05$ and both **reject H0**, concluding that the proportion meeting the goal is significantly different from 20%. The observed proportion is 22.53%, slightly above the hypothesised 20%.

The two p-values are very close (chi-squared: 0.0142; binomial: 0.0155) but not identical:

- The **chi-squared test** uses a large-sample normal approximation (valid here since $np_0 = 300$ and $n(1-p_0) = 1200$ are both $\gg 5$). It is computationally simpler and equivalent to a squared z-test.
- The **binomial test** is exact – it computes probabilities directly from the $\text{Binomial}(1500, 0.20)$ distribution without any approximation. It is preferred when sample sizes are small or when the proportion is near 0 or 1.
- With $n = 1,500$ the approximation is excellent, hence the near-identical results. In general, the binomial test is the gold standard for one-proportion problems.

Summary of All Results

Table 8: Complete summary of hypothesis tests (Topics G and H, $\alpha = 0.05$)

Topic	Question	Test	p-value	Decision
G	Fitness score before = 50?	One-sample t-test	0.081	Fail to reject H0
G	Mean weight loss != 5 kg?	One-sample t-test	< 0.001	Reject H0
G	Fitness improved after 8 weeks?	Paired t-test	< 0.001	Reject H0
G	Cholesterol change != 0?	t-test + Wilcoxon	< 0.001	Reject H0
G	RT before = 275 ms? (Male)	One-sample t-test	< 0.001	Reject H0
G	RT before = 275 ms? (Female)	One-sample t-test	< 0.001	Reject H0
G	RT improved? (whole cohort)	One-sample t-test	< 0.001	Reject H0
H	Male satisfaction uniform?	Chi-squared GoF	0.155	Fail to reject H0
H	≥ 90 pct Low Carb men lose weight?	z-test (one proportion)	0.0008	Reject H0
H	Prop. meeting goal = 20 pct?	Chi-squared + Binomial	0.014 / 0.016	Reject H0

Appendix: Session Info

```
sessionInfo()
```

```
## R version 4.5.2 (2025-10-31)
## Platform: aarch64-apple-darwin20
## Running under: macOS Tahoe 26.4.1
##
## Matrix products: default
## BLAS: /System/Library/Frameworks/Accelerate.framework/Versions/A/Frameworks/vecLib.framework/Versions/A
## LAPACK: /Library/Frameworks/R.framework/Versions/4.5-arm64/Resources/lib/libRlapack.dylib; LAPACK v
##
```

```

## locale:
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
##
## time zone: Australia/Sydney
## tzcode source: internal
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] kableExtra_1.4.0 knitr_1.51      lubridate_1.9.5  forcats_1.0.1
## [5] stringr_1.6.0    dplyr_1.2.0     purrr_1.2.1     readr_2.2.0
## [9] tidyr_1.3.2      tibble_3.3.1    ggplot2_4.0.2    tidyverse_2.0.0
##
## loaded via a namespace (and not attached):
## [1] gtable_0.3.6      compiler_4.5.2    tidyselect_1.2.1  xml2_1.5.2
## [5] textshaping_1.0.5 systemfonts_1.3.2 scales_1.4.0      yaml_2.3.12
## [9] fastmap_1.2.0     R6_2.6.1          labeling_0.4.3    generics_0.1.4
## [13] svglite_2.2.2     pillar_1.11.1     RColorBrewer_1.1-3 tzdb_0.5.0
## [17] rlang_1.1.7       stringi_1.8.7     xfun_0.56         S7_0.2.1
## [21] viridisLite_0.4.3 timechange_0.4.0  cli_3.6.5         withr_3.0.2
## [25] magrittr_2.0.4    digest_0.6.39     grid_4.5.2        rstudioapi_0.18.0
## [29] hms_1.1.4         lifecycle_1.0.5   vctrs_0.7.1       evaluate_1.0.5
## [33] glue_1.8.0        farver_2.1.2      rmarkdown_2.30    tools_4.5.2
## [37] pkgconfig_2.0.3   htmltools_0.5.9

```